

Spark Basics & Python

project with the latest technology

Lee Hae Joon



Reviews

What is RDD and Strong Points ?



Python

INSTALL JUPYTER NOTEBOOK

<https://www.digitalocean.com/community/tutorials/how-to-set-up-a-jupyter-notebook-to-run-ipython-on-ubuntu-16-04>



Python

Python is a high-level, interpreted, interactive and object-oriented scripting language. Python is designed to be highly readable.

1. **Interpreted:** Python is processed at runtime by the interpreter. You do not need to compile your program before executing it.
2. **Interactive:** You can actually sit at a Python prompt and interact with the interpreter directly to write your programs
3. **Object-Oriented:** Python supports Object-Oriented style or technique of programming that encapsulates code within objects.
4. **Beginner's Language:** Python is a great language for the beginner-level programmers and supports the development of a wide range of applications from simple text processing to WWW browsers to games.



Python

Reserved Keywords

and	exec	not
assert	finally	or
break	for	pass
class	from	print
continue	global	raise
def	if	return
del	import	try
elif	in	while
else	is	with
except	lambda	yield



Python Datatype



list - Like almost everything in Python, lists are objects. There are many methods associated to them. Some of which are presented here below.

```
In [1]: list_ = [1,2,3,'a','b','c']  
list_[0]
```

```
Out[1]: 1
```

```
In [2]: list_.append('d')  
list_
```

```
Out[2]: [1, 2, 3, 'a', 'b', 'c', 'd']
```

```
In [3]: list_append = [4,5,6]  
list_.append(list_append)  
list_
```

```
Out[3]: [1, 2, 3, 'a', 'b', 'c', 'd', [4, 5, 6]]
```

```
In [4]: list_extend = [7,8,9]  
list_.extend(list_extend)  
list_
```

```
Out[4]: [1, 2, 3, 'a', 'b', 'c', 'd', [4, 5, 6], 7, 8, 9]
```

```
In [5]: list_.index('b')
```

```
Out[5]: 4
```



```
In [5]: list_.index('b')
```

```
Out[5]: 4
```

```
In [7]: list_.insert(1, 'c')  
list_
```

```
Out[7]: [1, 'c', 'c', 2, 3, 'a', 'b', 'c', 'd', [4, 5, 6], 7, 8, 9]
```

```
In [8]: ## slicing  
list_[2:4]
```

```
Out[8]: ['c', 2]
```

```
In [46]: list_.sort()  
list_
```

```
Out[46]: [1, 2, 3, 7, 8, 9, [4, 5, 6], 'a', 'b', 'c', 'c', 'c', 'd']
```

```
In [10]: ## list comprehension  
evens = []  
for i in list_:  
    if list_.index(i)%2 == 0:  
        evens.append(i)  
  
evens
```

```
Out[10]: [1, 3, 'b', 'd', 7, 9]
```



Tuples - This is a data structure very similar to the list data structure. The main difference being that tuple manipulation are faster than list because tuples are immutable.

```
In [11]: ### Constructing tuples  
tuple_ = (1,2,3)  
tuple_[0]
```

```
Out[11]: 1
```

```
In [13]: tuple_[0] = 4  
tuple_
```

```
-----  
TypeError                                Traceback (most recent call last)  
<ipython-input-13-a7fa442c0cba> in <module>()  
----> 1 tuple_[0] = 4  
      2 tuple_
```

```
TypeError: 'tuple' object does not support item assignment
```



```
In [14]: tuple_2 = 4,5,6 # constructed by using comma  
tuple_2
```

```
Out[14]: (4, 5, 6)
```

```
In [16]: tuple_3 = tuple_2 * 3  
tuple_3
```

```
Out[16]: (4, 5, 6, 4, 5, 6, 4, 5, 6)
```

```
In [17]: ### interaction with dictionary  
dict_ = dict([('one',1), ('two',2), ('three',3)])  
dict_
```

```
Out[17]: {'one': 1, 'three': 3, 'two': 2}
```



Dictionary

```
In [23]: dict_ = {'one':1, 'two':2, 'three':3}
dict_.keys()
dict_.values()
```

```
Out[23]: [3, 2, 1]
```

```
In [24]: # iterating from dictionary
for ele in dict_.items():
    print ele
```

```
('three', 3)
('two', 2)
('one', 1)
```

```
In [26]: dict_.has_key('one')
```

```
Out[26]: True
```

```
In [27]: dict_.get('one')
```

```
Out[27]: 1
```

```
In [27]: dict_.get('one')
```

```
Out[27]: 1
```

```
In [28]: [x for x in dict_.itervalues()] # iterkeys(), iteritems()
```

```
Out[28]: [3, 2, 1]
```

```
In [33]: ### combining dictionaries  
dict_1 = {'a':1}  
dict_2 = {'b':2}  
dict_1.update(dict_2)  
dict_1
```

```
Out[33]: {'a': 1, 'b': 2}
```

```
In [34]: # deleting item from dictionary  
del dict_1['a']  
dict_1
```

```
Out[34]: {'b': 2}
```



Sets, constructed from a sequence (or some other iterable object). Since sets cannot have duplicated, there are usually used to build sequence of unique items (e.g., set of identifiers).

```
In [35]: a = set([1, 2, 3, 4])  
        b = set([3, 4, 5, 6])
```

```
In [37]: # union  
        a | b  
        # intersection  
        a & b
```

```
Out[37]: {3, 4}
```

```
In [39]: # difference  
        a - b
```

```
Out[39]: {1, 2}
```



```
In [39]: # difference  
a - b
```

```
Out[39]: {1, 2}
```

```
In [40]: # symmetric difference  
a ^ b
```

```
Out[40]: {1, 2, 5, 6}
```

```
In [44]: c = a.intersection(b) # c = a&b  
c
```

```
Out[44]: {3, 4}
```



DataFrame, Pandas

Next to Matplotlib and NumPy, [Pandas](#) is one of the most widely used Python libraries in data science. It is mainly used for data munging, and with good reason: it's very powerful and flexible, among many other things. It makes the least sexy part of the "sexiest job of the 21st Century" a bit more pleasant.



```
In [11]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: s = pd.Series([1,3,5,np.nan,6,8])
```

```
In [12]: dates = pd.date_range('20170101', periods = 6)
```

```
In [8]: dates
```

```
Out[8]: <class 'pandas.tseries.index.DatetimeIndex'>
[2017-01-01, ..., 2017-01-06]
Length: 6, Freq: D, Timezone: None
```

```
In [ ]:
```

```
In [13]: df = pd.DataFrame(np.random.randn(6,4), index = dates, columns = list('ABCD'))
```

```
In [10]: df
```

```
Out[10]:
```

	A	B	C	D
2017-01-01	-0.671602	-1.199595	1.250647	-1.291698
2017-01-02	-0.837777	-1.087355	-1.751561	0.098106
2017-01-03	-0.874211	-0.671499	-0.059741	-1.065709
2017-01-04	1.860088	-0.056245	0.469709	0.226871
2017-01-05	-0.262813	0.279909	1.087980	-0.080228
2017-01-06	-0.901976	-1.081591	-1.451922	-1.438884

6 rows x 4 columns




```
In [13]: df.head(1)
```

```
Out[13]:
```

	A	B	C	D
2017-01-01	-0.671602	-1.199595	1.250647	-1.291698

1 rows x 4 columns

```
In [12]: df.tail(2)
```

```
Out[12]:
```

	A	B	C	D
2017-01-05	-0.262813	0.279909	1.087980	-0.080228
2017-01-06	-0.901976	-1.081591	-1.451922	-1.438884

2 rows x 4 columns

```
In [14]: df.describe()
```

```
Out[14]:
```

	A	B	C	D
count	6.000000	6.000000	6.000000	6.000000
mean	-0.281382	-0.636063	-0.075815	-0.591924
std	1.075608	0.615821	1.273858	0.753644
min	-0.901976	-1.199595	-1.751561	-1.438884
25%	-0.865103	-1.085914	-1.103876	-1.235201
50%	-0.754690	-0.876545	0.204984	-0.572969
75%	-0.365011	-0.210058	0.933412	0.053523
max	1.860088	0.279909	1.250647	0.226871

8 rows x 4 columns



```
In [15]: df.T
```

```
Out[15]:
```

	2017-01-01 00:00:00	2017-01-02 00:00:00	2017-01-03 00:00:00	2017-01-04 00:00:00	2017-01-05 00:00:00	2017-01-06 00:00:00
A	-0.671602	-0.837777	-0.874211	1.860088	-0.262813	-0.901976
B	-1.199595	-1.087355	-0.671499	-0.056245	0.279909	-1.081591
C	1.250647	-1.751561	-0.059741	0.469709	1.087980	-1.451922
D	-1.291698	0.098106	-1.065709	0.226871	-0.080228	-1.438884

4 rows x 6 columns

```
In [52]: df.sort_index(axis=0, ascending = False)
#df.sort_index(axis=1, ascending = False)
```

```
Out[52]:
```

	A	B	C	D
2017-01-06	-0.901976	-1.081591	-1.451922	-1.438884
2017-01-05	-0.262813	0.279909	1.087980	-0.080228
2017-01-04	1.860088	-0.056245	0.469709	0.226871
2017-01-03	-0.874211	-0.671499	-0.059741	-1.065709
2017-01-02	-0.837777	-1.087355	-1.751561	0.098106
2017-01-01	-0.671602	-1.199595	1.250647	-1.291698

6 rows x 4 columns

```
In [21]: df[:3] # selected rows label
```

```
Out[21]:
```

	A	B	C	D
2017-01-01	-0.671602	-1.199595	1.250647	-1.291698
2017-01-02	-0.837777	-1.087355	-1.751561	0.098106
2017-01-03	-0.874211	-0.671499	-0.059741	-1.065709

3 rows x 4 columns



```
In [21]: df[:3] # selected rows label
```

```
Out[21]:
```

	A	B	C	D
2017-01-01	-0.671602	-1.199595	1.250647	-1.291698
2017-01-02	-0.837777	-1.087355	-1.751561	0.098106
2017-01-03	-0.874211	-0.671499	-0.059741	-1.065709

3 rows × 4 columns

```
In [27]: df.loc['2017-01-01'] # label  
df.loc[dates[0]] # select by label 'row'
```

```
Out[27]: A    -0.671602  
        B    -1.199595  
        C     1.250647  
        D    -1.291698  
        Name: 2017-01-01 00:00:00, dtype: float64
```

```
In [32]: df.loc[:, 'A':'C'] # row label, column label  
df.loc[:, ['A', 'C']]
```

```
Out[32]:
```

	A	C
2017-01-01	-0.671602	1.250647
2017-01-02	-0.837777	-1.751561
2017-01-03	-0.874211	-0.059741
2017-01-04	1.860088	0.469709
2017-01-05	-0.262813	1.087980
2017-01-06	-0.901976	-1.451922

6 rows × 2 columns



```
In [34]: # selection by position
df.iloc[3]
df.iloc[3:5,0:2]
```

```
Out[34]:
```

	A	B
2017-01-04	1.860088	-0.056245
2017-01-05	-0.262813	0.279909

2 rows x 2 columns

```
In [37]: df_filter = df[df>0]
df_filter.fillna('zero') # df_filter.dropna(how='any')
```

```
Out[37]:
```

	A	B	C	D
2017-01-01	zero	zero	1.250647	zero
2017-01-02	zero	zero	zero	0.09810621
2017-01-03	zero	zero	zero	zero
2017-01-04	1.860088	zero	0.4697087	0.2268709
2017-01-05	zero	0.2799088	1.08798	zero
2017-01-06	zero	zero	zero	zero

6 rows x 4 columns

```
In [14]: df_copy = df.copy()
df_copy['copy'] = ['one', 'two', 'three', np.nan, np.nan, 'six']
df_copy
```

```
Out[14]:
```

	A	B	C	D	copy
2017-01-01	1.617212	-2.285564	-0.132234	0.273070	one
2017-01-02	-0.352639	0.357600	-0.571745	0.515890	two
2017-01-03	-1.097020	2.993803	-0.682899	-0.272870	three
2017-01-04	-0.296105	-1.641046	0.571109	-0.245860	NaN
2017-01-05	-0.672826	-0.015981	0.001862	-1.756203	NaN
2017-01-06	0.157500	-2.489007	-0.395584	-1.312094	six



```
In [15]: df_copy['A+10'] = df_copy['A'].apply(lambda x : x + 10)
df_copy
```

```
Out[15]:
```

	A	B	C	D	copy	A+10
2017-01-01	1.617212	-2.285564	-0.132234	0.273070	one	11.617212
2017-01-02	-0.352639	0.357600	-0.571745	0.515890	two	9.647361
2017-01-03	-1.097020	2.993803	-0.682899	-0.272870	three	8.902980
2017-01-04	-0.296105	-1.641046	0.571109	-0.245860	NaN	9.703895
2017-01-05	-0.672826	-0.015981	0.001862	-1.756203	NaN	9.327174
2017-01-06	0.157500	-2.489007	-0.395584	-1.312094	six	10.157500

```
In [44]: s = pd.Series([1,3,5,np.nan,6,8], index=dates).shift(2)
s
```

```
Out[44]:
```

2017-01-01	NaN
2017-01-02	NaN
2017-01-03	1
2017-01-04	3
2017-01-05	5
2017-01-06	NaN

Freq: D, dtype: float64

```
In [51]: df.sub(s,axis = 0) # axis = 'index'
```

```
Out[51]:
```

	A	B	C	D
2017-01-01	NaN	NaN	NaN	NaN
2017-01-02	NaN	NaN	NaN	NaN
2017-01-03	-1.874211	-1.671499	-1.059741	-2.065709
2017-01-04	-1.139912	-3.056245	-2.530291	-2.773129
2017-01-05	-5.262813	-4.720091	-3.912020	-5.080228
2017-01-06	NaN	NaN	NaN	NaN

6 rows x 4 columns



```
In [60]: df.apply(np.cumsum, axis=1)
df.apply(np.cumsum, axis=0)
```

```
Out[60]:
```

	A	B	C	D
2017-01-01	-0.671602	-1.199595	1.250647	-1.291698
2017-01-02	-1.509379	-2.286950	-0.500913	-1.193592
2017-01-03	-2.383590	-2.958449	-0.560654	-2.259301
2017-01-04	-0.523503	-3.014694	-0.090946	-2.032430
2017-01-05	-0.786316	-2.734785	0.997034	-2.112658
2017-01-06	-1.688292	-3.816376	-0.454887	-3.551542

6 rows x 4 columns

```
In [61]: df.apply(lambda x : x.max() - x.min())
```

```
Out[61]: A    2.762064
B    1.479504
C    3.002208
D    1.665755
dtype: float64
```



```
In [77]: df = pd.DataFrame({'A': ['one', 'one', 'two', 'three'] * 3, 'B': ['A', 'B', 'C'] * 4,
                          'C': ['foo', 'foo', 'foo', 'bar', 'bar', 'bar'] * 2, 'D': np.random.randn(12),
                          'E': np.random.randn(12)})
df
```

```
Out[77]:
```

		A	B	C	D	E
0	one	A	foo	0.453737	0.836380	
1	one	B	foo	0.281438	-2.513328	
2	two	C	foo	0.796836	1.051851	
3	three	A	bar	-0.610079	0.423231	
4	one	B	bar	-1.447885	0.359370	
5	one	C	bar	1.511768	-1.312828	
6	two	A	foo	1.102993	0.990530	
7	three	B	foo	-0.238541	-1.701424	
8	one	C	foo	-0.015979	0.922675	
9	one	A	bar	0.228019	-2.041269	
10	two	B	bar	2.906439	0.038860	
11	three	C	bar	-1.683287	0.361368	

12 rows x 5 columns

```
In [17]: df_val = pd.pivot_table(df, values='D', index=['A', 'B'], columns=['C'])
df_val
```

```
Out[17]:
```

		A	B	C	-0.682899448303	-0.571744507826	-0.395584066434	-0.132233839532	0.00186176615742	0.571109205329
	-1.097020	2.993803	-0.27287	NaN	NaN	NaN	NaN	NaN	NaN	NaN
	-0.672826	-0.015981	NaN	NaN	NaN	NaN	NaN	-1.756203	NaN	NaN
	-0.352639	0.357600	NaN	0.51589	NaN	NaN	NaN	NaN	NaN	NaN
	-0.296105	-1.641046	NaN	NaN	NaN	NaN	NaN	NaN	NaN	-0.24586
	0.157500	-2.489007	NaN	NaN	-1.312094	NaN	NaN	NaN	NaN	NaN
	1.617212	-2.285564	NaN	NaN	NaN	NaN	0.27307	NaN	NaN	NaN



Advanced DataFrame Pandas




```
In [3]: # Create a dataframe with dates as your index
States = ['NY', 'NY', 'NY', 'NY', 'FL', 'FL', 'GA', 'GA', 'FL', 'FL']
data = [1.0, 2, 3, 4, 5, 6, 7, 8, 9, 10]
idx = pd.date_range('1/1/2012', periods=10, freq='MS')
dfl = pd.DataFrame(data, index=idx, columns=['Revenue'])
dfl['State'] = States
dfl
```

Out[3]:

	Revenue	State
2012-01-01	1.0	NY
2012-02-01	2.0	NY
2012-03-01	3.0	NY
2012-04-01	4.0	NY
2012-05-01	5.0	FL
2012-06-01	6.0	FL
2012-07-01	7.0	GA
2012-08-01	8.0	GA
2012-09-01	9.0	FL
2012-10-01	10.0	FL

<https://pandas.pydata.org/pandas-docs/stable/timeseries.html>



```
In [4]: # Create a second dataframe
data2 = [10.0, 10.0, 9, 9, 8, 8, 7, 7, 6, 6]
idx2 = pd.date_range('1/1/2013', periods=10, freq='MS')
df2 = pd.DataFrame(data2, index=idx2, columns=['Revenue'])
df2['State'] = States
df2
```

```
Out[4]:
```

	Revenue	State
2013-01-01	10.0	NY
2013-02-01	10.0	NY
2013-03-01	9.0	NY
2013-04-01	9.0	NY
2013-05-01	8.0	FL
2013-06-01	8.0	FL
2013-07-01	7.0	GA
2013-08-01	7.0	GA
2013-09-01	6.0	FL
2013-10-01	6.0	FL



```
In [5]: df = pd.concat([df1,df2])
df
```

Out[5]:

	Revenue	State
2012-01-01	1.0	NY
2012-02-01	2.0	NY
2012-03-01	3.0	NY
2012-04-01	4.0	NY
2012-05-01	5.0	FL
2012-06-01	6.0	FL
2012-07-01	7.0	GA
2012-08-01	8.0	GA
2012-09-01	9.0	FL
2012-10-01	10.0	FL
2013-01-01	10.0	NY
2013-02-01	10.0	NY
2013-03-01	9.0	NY
2013-04-01	9.0	NY
2013-05-01	8.0	FL
2013-06-01	8.0	FL
2013-07-01	7.0	GA
2013-08-01	7.0	GA
2013-09-01	6.0	FL
2013-10-01	6.0	FL

	Revenue	State
2012-01-01	1.0	NY
2012-02-01	2.0	NY
2012-03-01	3.0	NY
2012-04-01	4.0	NY
2012-05-01	5.0	FL
2012-06-01	6.0	FL
2012-07-01	7.0	GA
2012-08-01	8.0	GA
2012-09-01	9.0	FL
2012-10-01	10.0	FL
2013-01-01	10.0	NY
2013-02-01	10.0	NY
2013-03-01	9.0	NY
2013-04-01	9.0	NY
2013-05-01	8.0	FL
2013-06-01	8.0	FL
2013-07-01	7.0	GA
2013-08-01	7.0	GA
2013-09-01	6.0	FL
2013-10-01	6.0	FL



```
In [7]: # Method 1

# make a copy of original df
newdf = df.copy()

newdf['x-Mean'] = abs(newdf['Revenue'] - newdf['Revenue'].mean())
newdf['1.96*std'] = 1.96*newdf['Revenue'].std()
newdf['Outlier'] = abs(newdf['Revenue'] - newdf['Revenue'].mean()) > 1.96*newdf['Revenue'].std()
newdf
```

```
Out[7]:
```

	Revenue	State	x-Mean	1.96*std	Outlier
2012-01-01	1.0	NY	5.75	5.200273	True
2012-02-01	2.0	NY	4.75	5.200273	False
2012-03-01	3.0	NY	3.75	5.200273	False
2012-04-01	4.0	NY	2.75	5.200273	False
2012-05-01	5.0	FL	1.75	5.200273	False
2012-06-01	6.0	FL	0.75	5.200273	False
2012-07-01	7.0	GA	0.25	5.200273	False
2012-08-01	8.0	GA	1.25	5.200273	False
2012-09-01	9.0	FL	2.25	5.200273	False
2012-10-01	10.0	FL	3.25	5.200273	False
2013-01-01	10.0	NY	3.25	5.200273	False
2013-02-01	10.0	NY	3.25	5.200273	False
2013-03-01	9.0	NY	2.25	5.200273	False
2013-04-01	9.0	NY	2.25	5.200273	False
2013-05-01	8.0	FL	1.25	5.200273	False
2013-06-01	8.0	FL	1.25	5.200273	False
2013-07-01	7.0	GA	0.25	5.200273	False
2013-08-01	7.0	GA	0.25	5.200273	False
2013-09-01	6.0	FL	0.75	5.200273	False
2013-10-01	6.0	FL	0.75	5.200273	False



```
In [26]: # Method 2
# Group by item

# make a copy of original df
newdf = df.copy()

State = newdf.groupby('State')
State

for idx, data in State: # DataFrameGroupBy
    print idx, data
```

```
FL          Revenue State
2012-05-01    5.0    FL
2012-06-01    6.0    FL
2012-09-01    9.0    FL
2012-10-01   10.0    FL
2013-05-01    8.0    FL
2013-06-01    8.0    FL
2013-09-01    6.0    FL
2013-10-01    6.0    FL
GA          Revenue State
2012-07-01    7.0    GA
2012-08-01    8.0    GA
2013-07-01    7.0    GA
2013-08-01    7.0    GA
NY          Revenue State
2012-01-01    1.0    NY
2012-02-01    2.0    NY
2012-03-01    3.0    NY
2012-04-01    4.0    NY
2013-01-01   10.0    NY
2013-02-01   10.0    NY
2013-03-01    9.0    NY
2013-04-01    9.0    NY
```



```

In [8]: # Method 2
        # Group by multiple items

        # make a copy of original df
        newdf = df.copy()

        StateMonth = newdf.groupby(['State', lambda x: x.month])

        newdf['Outlier'] = StateMonth.transform( lambda x: abs(x-x.mean()) > 1.96*x.std() )
        newdf['x-Mean'] = StateMonth.transform( lambda x: abs(x-x.mean()) )
        newdf['1.96*std'] = StateMonth.transform( lambda x: 1.96*x.std() )
        newdf

```

```

Out[8]:

```

	Revenue	State	Outlier	x-Mean	1.96*std
2012-01-01	1	NY	False	4.5	12.473364
2012-02-01	2	NY	False	4.0	11.087434
2012-03-01	3	NY	False	3.0	8.315576
2012-04-01	4	NY	False	2.5	6.929646
2012-05-01	5	FL	False	1.5	4.157788
2012-06-01	6	FL	False	1.0	2.771859
2012-07-01	7	GA	False	0.0	0.000000
2012-08-01	8	GA	False	0.5	1.385929
2012-09-01	9	FL	False	1.5	4.157788
2012-10-01	10	FL	False	2.0	5.543717
2013-01-01	10	NY	False	4.5	12.473364
2013-02-01	10	NY	False	4.0	11.087434
2013-03-01	9	NY	False	3.0	8.315576
2013-04-01	9	NY	False	2.5	6.929646
2013-05-01	8	FL	False	1.5	4.157788
2013-06-01	8	FL	False	1.0	2.771859
2013-07-01	7	GA	False	0.0	0.000000
2013-08-01	7	GA	False	0.5	1.385929
2013-09-01	6	FL	False	1.5	4.157788
2013-10-01	6	FL	False	2.0	5.543717

20 rows x 5 columns



```

In [9]: # Method 3
        # Group by item

        # make a copy of original df
        newdf = df.copy()

        State = newdf.groupby('State')

        def s(group):
            group['x-Mean'] = abs(group['Revenue'] - group['Revenue'].mean())
            group['1.96*std'] = 1.96*group['Revenue'].std()
            group['Outlier'] = abs(group['Revenue'] - group['Revenue'].mean()) > 1.96*group['Revenue'].std()
            return group

        Newdf2 = State.apply(s)
        Newdf2

```

```

Out[9]:

```

	Revenue	State	x-Mean	1.96*std	Outlier
2012-01-01	1	NY	5.00	7.554813	False
2012-02-01	2	NY	4.00	7.554813	False
2012-03-01	3	NY	3.00	7.554813	False
2012-04-01	4	NY	2.00	7.554813	False
2012-05-01	5	FL	2.25	3.434996	False
2012-06-01	6	FL	1.25	3.434996	False
2012-07-01	7	GA	0.25	0.980000	False
2012-08-01	8	GA	0.75	0.980000	False
2012-09-01	9	FL	1.75	3.434996	False
2012-10-01	10	FL	2.75	3.434996	False
2013-01-01	10	NY	4.00	7.554813	False
2013-02-01	10	NY	4.00	7.554813	False
2013-03-01	9	NY	3.00	7.554813	False
2013-04-01	9	NY	3.00	7.554813	False
2013-05-01	8	FL	0.75	3.434996	False
2013-06-01	8	FL	0.75	3.434996	False
2013-07-01	7	GA	0.25	0.980000	False
2013-08-01	7	GA	0.25	0.980000	False
2013-09-01	6	FL	1.25	3.434996	False



```
In [8]: file_name = "./analyzed_american.csv"  
Newdf2.to_csv(file_name)
```

☐  analyzed_american.csv

2 minutes ago



Matplotlib



```
In [6]: ts = pd.Series(np.random.randn(1000), index=pd.date_range('1/1/2000', periods=1000))
ts
```


```
Out[6]: 2000-01-01    1.075219
2000-01-02    0.436073
2000-01-03   -0.901465
2000-01-04   -0.133734
2000-01-05   -1.400580
2000-01-06    1.506402
2000-01-07    2.166358
2000-01-08    0.526718
2000-01-09    0.148402
2000-01-10   -1.341403
-----
```

```
In [1]: %matplotlib inline
```

```
In [16]: ts = ts.cumsum()
ts.plot()
ts.hist()

#matplotlib reference - https://matplotlib.org/tutorials/introductory/sample\_plots.html#sphx-glr-tutorials-introductory-sample-plots-py
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x113312d50>
```



https://matplotlib.org/tutorials/introductory/sample_plots.html#sphx-glr-tutorials-introductory-sample-plots-py



Practice

room_id	host_id	room_type	borough	neighborhood	reviews	overall_satisfaction	accommodates	bedrooms	price	minstay	latitude	longitude	last_modified
6291807	16375951	Entire home/apt		Jamaica Plain	1		4	2	119	14	42.29816	-71.11153	2016-05-19 02:58:16.563871
2656568	13597630	Entire home/apt		Back Bay	0		2	1	600	1	42.348072	-71.076639	2016-05-19 02:58:06.015822
10723203	15913699	Private room		Allston	2		2	1	96	1	42.350588	-71.129477	2016-05-19 02:57:39.074104
10034592	20399668	Private room		Dorchester	13	5	2	1	55	1	42.317168	-71.040483	2016-05-19 02:57:28.669274
5454513	4962900	Entire home/apt		Back Bay	13	4.5	5	2	276		42.346598	-71.080123	2016-05-19 02:56:56.182103
335730	290698	Private room		East Boston	34	3.5	9	1	70	1	42.39001	-70.996161	2016-05-19 02:56:44.188597
7635616	22348222	Entire home/apt		Beacon Hill	3	5	2	0	155	7	42.359601	-71.063908	2016-05-19 02:56:18.932562
12808014	28197086	Private room		Allston	0		1	1	65	1	42.351883	-71.130772	2016-05-19 02:56:12.350833
6793913	30283594	Entire home/apt		Fenway	2		5	2	571	3	42.343305	-71.101489	2016-05-19 02:55:56.997256
586994	2894162	Entire home/apt		Beacon Hill	47	4.5	2	1	160	4	42.358926	-71.063109	2016-05-19 02:54:56.750137



```
In [4]: import pandas as pd
import numpy as np
```

```
In [5]: basic_folder = '/Users/phil/python_work/venv/'
file_name = basic_folder + 'airbnb.csv'
df_air = pd.read_csv(file_name)
```

```
In [6]: df_air
```

```
Out[6]:
```

	room_id	host_id	room_type	borough	neighborhood	reviews	overall_satisfaction	accommodates	bedrooms	price	minstay	latitude	longitude
0	6291807	16375951	Entire home/apt	NaN	Jamaica Plain	1	NaN	4	2.0	119.0	14.0	42.298160	-71.111530
1	2656568	13597630	Entire home/apt	NaN	Back Bay	0	NaN	2	1.0	600.0	1.0	42.348072	-71.076639
2	10723203	15913699	Private room	NaN	Allston	2	NaN	2	1.0	96.0	1.0	42.350588	-71.129477
3	10034592	20399668	Private room	NaN	Dorchester	13	5.0	2	1.0	55.0	1.0	42.317168	-71.040483
4	5454513	4962900	Entire home/apt	NaN	Back Bay	13	4.5	5	2.0	276.0	NaN	42.346598	-71.080123
5	335730	290698	Private room	NaN	East Boston	34	3.5	9	1.0	70.0	1.0	42.390010	-70.996161
6	7635616	22348222	Entire home/apt	NaN	Beacon Hill	3	5.0	2	0.0	155.0	7.0	42.359601	-71.063908

```
In [8]: df_air['overall_satisfaction'] = df_air.overall_satisfaction.fillna(3)
```

```
In [9]: df_air
```

3270	6248970	32449720	Private room	NaN	Jamaica Plain	52	5.0	1	1.0	65.0	3.0	42.321720	-71.103459
3271	7549896	39568178	Private room	NaN	Chinatown	3	5.0	2	1.0	100.0	7.0	42.347990	-71.064152
3272	8519646	15098486	Private room	NaN	North End	16	4.5	2	1.0	116.0	1.0	42.364245	-71.052945
3273	6574771	16881770	Private room	NaN	Dorchester	9	5.0	1	1.0	40.0	3.0	42.323962	-71.058513
3274	1321422	6608084	Private room	NaN	Dorchester	168	5.0	3	1.0	45.0	1.0	42.308380	-71.046943
3275	12590656	51449558	Private room	NaN	Dorchester	0	3.0	3	1.0	55.0	2.0	42.321141	-71.056032
3276	1178371	6430732	Private room	NaN	Dorchester	66	5.0	1	1.0	45.0	2.0	42.294568	-71.066813

3277 rows x 14 columns

```
In [10]: df_neigh = df_air.set_index(['neighborhood'])
df_neigh
```

```
Out[10]:
```

neighborhood	room_id	host_id	room_type	borough	reviews	overall_satisfaction	accommodates	bedrooms	price	minstay	latitude	longitude	last_review
Jamaica Plain	6291807	16375951	Entire home/apt	NaN	1	3.0	4	2.0	119.0	14.0	42.298160	-71.111530	2015-02-28
Back Bay	2656568	13597630	Entire home/apt	NaN	0	3.0	2	1.0	600.0	1.0	42.348072	-71.076639	2015-02-28
Allston	10723203	15913699	Private room	NaN	2	3.0	2	1.0	96.0	1.0	42.350588	-71.129477	2015-02-27
Dorchester	10034592	20399668	Private room	NaN	13	5.0	2	1.0	55.0	1.0	42.317168	-71.040483	2015-02-27
Back Bay	5454513	4962900	Entire home/apt	NaN	13	4.5	5	2.0	276.0	NaN	42.346598	-71.080123	2015-02-26
East Boston	335730	290698	Private room	NaN	34	3.5	9	1.0	70.0	1.0	42.390010	-70.996161	2015-02-26
Beacon Hill	7635616	22348222	Entire	NaN	3	5.0	2	0.0	155.0	7.0	42.359601	-71.063908	2015-02-26



Answer 4 questions

